Copy page

Metrics API

```
GET /api/public/metrics
```

The **Metrics API** enables you to retrieve customized analytics from your Langfuse data. This endpoint allows you to specify dimensions, metrics, filters, and time granularity to build powerful custom reports and dashboards for your LLM applications.

Metrics API v2 (Beta)

⚠️ The v2 Metrics API is currently in **beta**. The API is stable for production use, but some parameters and behaviors may change based on user feedback before general availability.

Cloud-only (Beta): The v2 Metrics API is only available on Langfuse Cloud and currently in beta. We are working on a robust migration path for self-hosted deployments.

Data availability note: When using current SDK versions, data may take approximately 5 minutes to appear on v2 endpoints. We will be releasing updated SDK versions soon that will make data available immediately.

```
GET /api/public/v2/metrics
```

The v2 Metrics API provides significant performance improvements through an optimized data architecture built on a new events table schema that minimizes database work per query.

Key Changes from v1

Ask AI

The `traces` view is no longer available in v2. Instead, use the `observations` view which is both faster and more powerful compared to v1.

Available Views in v2

View	Description
`observations`	Query observation-level data with optional trace-level aggregations
`scores-numeric`	Query numeric and boolean scores
`scores-categorical`	Query categorical (string) scores

Row Limit

The v2 Metrics API enforces a default `rowLimit` of 100 rows per query to ensure consistent performance. You can specify a custom `rowLimit` in your query to override this default.

High Cardinality Dimensions

Certain dimensions like `id`, `traceId`, `userId`, and `sessionId` cannot be used for grouping in the v2 Metrics API. Grouping by these high cardinality fields is extremely expensive and rarely useful in practice. These dimensions remain available for filtering.

Example: Most expensive models used in observations

```
curl \
  -H "Authorization: Basic <BASIC AUTH HEADER>" \
  -G \
  --data-urlencode 'query={
    "view": "observations",
    "metrics": [{"measure": "totalCost", "aggregation": "sum"}],
    "dimensions": [{"field": "providedModelName"}],
    "filters": [],
    "fromTimestamp": "2025-12-01T00:00:00Z",
    "toTimestamp": "2025-12-16T00:00:00Z",
```

```
    "orderBy": [{"field": "totalCost_sum", "direction": "desc"}]
  }' \
  https://cloud.langfuse.com/api/public/v2/metrics
```

API Reference: See the full v2 Metrics API Reference for all available parameters, response schemas, and interactive examples.

Metrics API v1

The Metrics API supports querying across different views (traces, observations, scores) and allows you to:

- Select specific dimensions to group your data
- Apply multiple metrics with different aggregation methods
- Filter data based on metadata, timestamps, and other properties
- Analyze data across time with customizable granularity
- Order results according to your needs

Query Parameters

The API accepts a JSON query object passed as a URL-encoded parameter:

Parameter	Type	Description
`query`	JSON string	The encoded query object defining what metrics to retrieve

Query Object Structure

Field	Type	Required	Description
`view`	string	Yes	The data view to query: `"traces"`, `"observations"`, `"scores-`

Field	Type	Required	Description
			`numeric"`, or `"scores-categorical"`
`dimensions`	array	No	Array of dimension objects to group by, e.g. `[{ "field": "name" }]`
`metrics`	array	Yes	Array of metric objects to calculate, e.g. `[{ "measure": "latency", "aggregation": "p95" }]`
`filters`	array	No	Array of filter objects to narrow results, e.g. `[{ "column": "metadata", "operator": "contains", "key": "customKey", "value": "customValue", "type": "stringObject" }]`
`timeDimension`	object	No	Configuration for time-based analysis, e.g. `{ "granularity": "day" }`
`fromTimestamp`	string	Yes	ISO timestamp for the start of the query period
`toTimestamp`	string	Yes	ISO timestamp for the end of the query period
`orderBy`	array	No	Specification for result ordering, e.g. `[{ "field": "name", "direction": "asc" }]`

Dimension Object Structure

```
{ "field": "name" }
```

Metric Object Structure

```
{ "measure": "count", "aggregation": "count" }
```

Common measure types include:

- `count` - Count of records
- `latency` - Duration/latency metrics

Aggregation types include:

- `sum` - Sum of values
- `avg` - Average of values
- `count` - Count of records
- `max` - Maximum value
- `min` - Minimum value
- `p50` - 50th percentile
- `p75` - 75th percentile
- `p90` - 90th percentile
- `p95` - 95th percentile
- `p99` - 99th percentile

Filter Object Structure

```json
{
  "column": "metadata",
  "operator": "contains",
  "key": "customKey",
  "value": "customValue",
  "type": "stringObject"
}
```

Time Dimension Object

```json
{
  "granularity": "day"
}
```

Supported granularities include: `hour`, `day`, `week`, `month`, and `auto`.

Example

Here's an example of querying the number of traces grouped by name:

API Python SDK

```
curl \
-H "Authorization: Basic <BASIC AUTH HEADER>" \
-G \
--data-urlencode 'query={
  "view": "traces",
  "metrics": [{"measure": "count", "aggregation": "count"}],
  "dimensions": [{"field": "name"}],
  "filters": [],
  "fromTimestamp": "2025-05-01T00:00:00Z",
  "toTimestamp": "2025-05-13T00:00:00Z"
}' \
https://cloud.langfuse.com/api/public/metrics
```

Response:


```
{
  "data": [
    { "name": "trace-test-2", "count_count": "10" },
    { "name": "trace-test-3", "count_count": "5" },
    { "name": "trace-test-1", "count_count": "3" }
  ]
}
```


Data Model

The Metrics API provides access to several data views, each with its own set of dimensions and metrics you can query. This section outlines the available options for each view.

Available Views

View	Description
`traces`	Query data at the trace level
`observations`	Query data at the observation level
`scores-numeric`	Query numeric and boolean scores
`scores-categorical`	Query categorical (string) scores

Trace Dimensions

Dimension	Type	Description
`id`	string	Trace ID
`name`	string	Trace name
`tags`	string[]	Trace tags
`userId`	string	User ID associated with the trace
`sessionId`	string	Session ID associated with the trace
`release`	string	Release tag
`version`	string	Version tag
`environment`	string	Environment (e.g., production, staging)
`observationName`	string	Name of related observations
`scoreName`	string	Name of related scores

Trace Metrics

Metric	Description
`count`	Count of traces
`observationsCount`	Count of observations within traces
`scoresCount`	Count of scores within traces
`latency`	Trace duration in milliseconds

Metric	Description
`totalTokens`	Total tokens used in the trace
`totalCost`	Total cost of the trace

Observation Dimensions

Dimension	Type	Description
`id`	string	Observation ID
`traceId`	string	Associated trace ID
`traceName`	string	Name of the parent trace
`environment`	string	Environment (e.g., production, staging)
`parentObservationId`	string	ID of parent observation
`type`	string	Observation type
`name`	string	Observation name
`level`	string	Log level
`version`	string	Version
`providedModelName`	string	Model name
`promptName`	string	Prompt name
`promptVersion`	string	Prompt version
`userId`	string	User ID from parent trace
`sessionId`	string	Session ID from parent trace
`traceRelease`	string	Release from parent trace
`traceVersion`	string	Version from parent trace
`scoreName`	string	Related score name

Observation Metrics

Metric	Description
`count`	Count of observations
`latency`	Observation duration in milliseconds
`totalTokens`	Total tokens used
`totalCost`	Total cost
`timeToFirstToken`	Time to first token in milliseconds
`countScores`	Count of related scores

Score Dimensions (Common)

Dimension	Type	Description
`id`	string	Score ID
`name`	string	Score name
`environment`	string	Environment
`source`	string	Score source
`dataType`	string	Data type
`traceId`	string	Related trace ID
`traceName`	string	Related trace name
`userId`	string	User ID from trace
`sessionId`	string	Session ID from trace
`observationId`	string	Related observation ID
`observationName`	string	Related observation name
`observationModelName`	string	Model used in related observation
`observationPromptName`	string	Prompt name used in related observation
`observationPromptVersion`	string	Prompt version used in related observation

Dimension	Type	Description
`configId`	string	Configuration ID

Score Metrics

Numeric Scores

Metric	Description
`count`	Count of scores
`value`	Numeric score value

Categorical Scores

Metric	Description
`count`	Count of scores

Categorical scores have an additional dimension:

Dimension	Type	Description
`stringValue`	string	String value of the categorical score

Daily Metrics API (Legacy)

⚠️ This API is a legacy API. For new use cases, please use the [Metrics API] instead. It has higher rate-limits and offers more flexibility.

```
GET /api/public/metrics/daily
```

Via the **Daily Metrics API**, you can retrieve aggregated daily usage and cost metrics from Langfuse for downstream use, e.g., in analytics, billing, and rate-limiting. The API allows you to filter by application type, user, or tags for tailored data retrieval.

See [API reference]() for more details.

Overview

Returned data includes daily timeseries of:

- [Cost]() in USD
- Trace and observation count
- Break down by model name
- Usage (e.g. tokens) broken down by input and output usage
- [Cost]() in USD
- Trace and observation count

Optional filters:

- `traceName` to commonly filter by the application type, depending on how you use `name` in your traces
- `userId` to filter by [user]()
- `tags` to filter by [tags]()
- `fromTimestamp`
- `toTimestamp`

Missing a key metric or filter? Request it via our [idea board]().

Example

```
GET /api/public/metrics/daily?traceName=my-copilot&userId=john&limit=2
```

```
{
  "data": [
    {
      "date": "2024-02-18",
```

```
      "countTraces": 1500,
      "countObservations": 3000,
      "totalCost": 102.19,
      "usage": [
        {
          "model": "llama2",
          "inputUsage": 1200,
          "outputUsage": 1300,
          "totalUsage": 2500,
          "countTraces": 1000,
          "countObservations": 2000,
          "totalCost": 50.19
        },
        {
          "model": "gpt-4",
          "inputUsage": 500,
          "outputUsage": 550,
          "totalUsage": 1050,
          "countTraces": 500,
          "countObservations": 1000,
          "totalCost": 52.0
        }
      ]
    },
    {
      "date": "2024-02-17",
      "countTraces": 1250,
      "countObservations": 2500,
      "totalCost": 250.0,
      "usage": [
        {
          "model": "llama2",
          "inputUsage": 1000,
          "outputUsage": 1100,
          "totalUsage": 2100,
          "countTraces": 1250,
          "countObservations": 2500,
          "totalCost": 250.0
        }
      ]
    }
  ],
  "meta": {
    "page": 1,
    "limit": 2,
    "totalItems": 60,
    "totalPages": 30
  }
}
```

Last updated on December 22, 2025

Was this page helpful?

Support

Platform

- LLM Tracing
- Prompt Management
- Evaluation
- Human Annotation
- Datasets
- Metrics
- Playground

Integrations

- Python SDK
- JS/TS SDK
- OpenAI SDK
- Langchain
- Llama-Index
- Litellm
- Dify
- Flowise
- Langflow
- Vercel AI SDK
- Instructor
- API

Resources

- Documentation
- Interactive Demo
- Video demo (10 min)
- Changelog
- Roadmap
- Pricing
- Enterprise
- Self-hosting
- Open Source
- Why Langfuse?
- AI Engineering Library

About

- Blog
- Careers
- About us
- Press
- Customers
- Support
- Talk to us
- OSS Friends
- Twitter
- LinkedIn

Status

🇯🇵 Japanese

🇰🇷 Korean

🇨🇳 Chinese

Legal

Security

Imprint

Terms

Privacy

SOC 2 Type II

ISO 27001

GDPR

HIPAA